


Kubernetes (K8s)




* The entire explanation will be based on 1 master 3 workers architecture but there is no rule in this but in dev environment 1 master multiple workers but in production master needs some backup/relicas.

(i) control plane

- (a) scheduler
- (b) etcd
- (c) API server
- (d) Control Manager
- (e) Kubeproxy
- (f) kubelet

(g) CRI
(Container Runtime Interface)

Kubernetes is a container orchestration distributed software, works using pods.



1. API server:-

Entry point to the cluster.

All communication goes through it.

Handles REST requests. (kubectl uses this)

2. etcd:-

Distributed key-value store

Stores cluster configuration and state.

3. scheduler:-

Assigns Pods to worker nodes.

Decides where workloads run.

4. control Manager:-

Ensures desired state matches actual state

Handles replication, node health etc.

5. Kubelet

Agent running on each node
communicates with control plane
Ensures containers are running.

6. kube proxy

Manages network rules.

Enables service communication.

7. Container Runtime

Runs container (e.g. Docker, containerd)
core Application concepts

(i) Pod

smallest ~~deployable~~ deployable unit
in Kubernetes.
can contain one or multiple containers.

Shares:

Network

storage

IP address

Pods are ephemeral (temporary)

(ii) Deployment

Manages Replicasets

Enables rolling updates

Ensures desired number of pods are running.

(iii) Service:-

Enforces pods to Network
provides stable IP & DNS name.

Types:-

- (i) ClusterIP
- (ii) NodePort
- (iii) LoadBalancer
- (iv) ExternalName.

Calico handles the entire routing and inter node communication. Pods are running on different ~~host~~ nodes so they need to be reachable via inter node communication and Calico handles that quite impressively.